

# Assignment # 5, MIE 1628 Cloud Data Analytics

Student Name: Nazib Chowdhury

Student ID: # 1004694235

---

## 1 Part A:

### 1.1 A.1:

1. Ingest -> Azure Data Factory
2. Data Store -> Azure Data Lake
3. Prepare and Transform Data -> Azure Databricks
4. Model and Serve Data -> Azure Synapse Analytics

### 1.2 A.2:

**Explain how Stream Analytics works in Azure. Mention at least two common use cases or applications for this service**

Azure Stream Analytics is tailored for real time streaming data analytics tasks. So it can consume data in real time, very quickly, and process it according to our given functions and rapidly send the results to databases, online dashboards or apps in real time.

#### **Workflow of Azure Stream Analytics**

1. Input from Sources:  
typically input from IoT devices, with ingestion triggered by Azure Event Hubs, and then transmitted to Azure Blob Storage
2. Processing of Data:  
The data is automatically processed with SQL like language to undergo some transformations that we want before storing it.
3. Storing and Analytics:  
The output data from the processing is finally sent to store in SQL Database, or sent for visualization using tools like Power BI for real time data analytics.

#### **Common Use cases of this application:**

1. **IoT Device Data:**  
IoT device data collection happens in real time and with a very fast frequency, often around (10-30 Hz) aligning with the frequency of the sensor. For more higher frequency physical processes like vibration, the data collection frequency can be even higher. For this reason, IoT device monitoring requires stream analytics to be done in real time. Because in the event of machine failure from say overheating, the symptoms arise very fast

and does not allow for much time before the system breaks down. Therefore, IoT device monitoring require real time analytics

2. **Real Time Fraud Detection:**

Fraud detection is another useful case of stream analytics. Fraudulent operations take place very fast, and hence these activities like fraudulent website cookies, pop-ups, emails etc and even fraudulent transaction requests need to be monitored in real time.

1.3 A.3

Microsoft Azure

Search resources, services, and docs (G+)

Copilot

nazib.chowdhury@mail...  
UNIVERSITY OF TORONTO (UTO...

All services > StreamAnalyticsJob | Overview > streamanalytics1

<> streamanalytics1 | Query

Stream Analytics job

Search

Start job Open in VS Code Diagnostic settings Refresh Query language docs Share feedback Tutorial Job ready to start

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Resource visualizer

Job topology

Inputs

Functions

<> Query

Outputs

No-code editor (preview)

Settings

Developer tools

Monitoring

Automation

Help

Inputs (1)

Outputs (1)

Functions (0)

Test query Save query Discard changes

1 SELECT \*  
2 INTO BlobOutput  
3 FROM iothubnazib1  
4 HAVING Temperature > 25

Input preview Test results Job simulation (preview)

Download results

messageId	deviceId	temperature	humidity	EventProcessedUtcTime
bigint	string	float	float	datetime
66	"Raspberry Pi Web Cli...	26.28883170573294	78.97814679257893	"2025-08-13T16:08:38....
68	"Raspberry Pi Web Cli...	29.3785480244885	76.50581458785743	"2025-08-13T16:08:38....
69	"Raspberry Pi Web Cli...	26.02164171583974	70.49353395335453	"2025-08-13T16:08:38....
70	"Raspberry Pi Web Cli...	26.630333672916322	77.51411030892282	"2025-08-13T16:08:38....
74	"Raspberry Pi Web Cli...	27.92066395014235	63.843855072732644	"2025-08-13T16:08:38....
75	"Raspberry Pi Web Cli...	25.88139320880864	71.64465296825117	"2025-08-13T16:08:38....
77	"Raspberry Pi Web Cli...	28.988049128550063	61.561696995866775	"2025-08-13T16:08:38....
78	"Raspberry Pi Web Cli...	27.885728926206333	74.36909604993168	"2025-08-13T16:08:38....
79	"Raspberry Pi Web Cli...	27.147185344821146	74.70196878399065	"2025-08-13T16:08:38....
80	"Raspberry Pi Web Cli...	26.752486351885324	63.12819125266602	"2025-08-13T16:08:38....
82	"Raspberry Pi Web Cli...	30.436258846362684	61.90572074649573	"2025-08-13T16:08:38....
83	"Raspberry Pi Web Cli...	26.29919647421592	63.08419765636294	"2025-08-13T16:08:38....
84	"Raspberry Pi Web Cli...	31.23902991143045	67.88082623429722	"2025-08-13T16:08:38....
86	"Raspberry Pi Web Cli...	29.656646523320767	68.16021948717815	"2025-08-13T16:08:38....
88	"Raspberry Pi Web Cli...	29.319218752933153	61.306406943565506	"2025-08-13T16:08:38....
92	"Raspberry Pi Web Cli...	26.47839850183596	66.18060256663826	"2025-08-13T16:08:38....
93	"Raspberry Pi Web Cli...	26.114092001711043	77.44131929370911	"2025-08-13T16:08:38....
94	"Raspberry Pi Web Cli...	25.55252873322818	60.3444223380991	"2025-08-13T16:08:38....
95	"Raspberry Pi Web Cli...	28.55107101458713	68.70672359724904	"2025-08-13T16:08:38....

Showing 31 rows from 'BlobOutput'. Ln 4, Col 24

Help

95	"Raspberry Pi Web Cli...	28.55107101458713	68.70672359724904	"2025-08-13T16:08:38....
97	"Raspberry Pi Web Cli...	27.03418722556435	78.25207434183129	"2025-08-13T16:08:38....
98	"Raspberry Pi Web Cli...	30.209301107878147	66.83359623173331	"2025-08-13T16:08:38....
99	"Raspberry Pi Web Cli...	28.386052721899613	71.50090093463587	"2025-08-13T16:08:38....
101	"Raspberry Pi Web Cli...	28.386356017363205	62.88035589377262	"2025-08-13T16:08:38....
102	"Raspberry Pi Web Cli...	26.648866322816282	68.54532277913447	"2025-08-13T16:08:38....
103	"Raspberry Pi Web Cli...	28.48288086763896	69.55304758386781	"2025-08-13T16:08:38....
104	"Raspberry Pi Web Cli...	28.087426169706983	76.54001249129192	"2025-08-13T16:08:38....
105	"Raspberry Pi Web Cli...	25.01510649151131	71.20054443953836	"2025-08-13T16:08:38....
107	"Raspberry Pi Web Cli...	31.85667664374832	72.28915083553807	"2025-08-13T16:08:38....
109	"Raspberry Pi Web Cli...	29.75480562095099	60.04594713068053	"2025-08-13T16:08:38....
111	"Raspberry Pi Web Cli...	30.39763178832444	72.44579529289089	"2025-08-13T16:08:38....
112	"Raspberry Pi Web Cli...	26.17712856434244	65.49371494288415	"2025-08-13T16:08:38....

Showing 31 rows from 'BlobOutput'.Ln 4, Col 24

2 Part B:

2.1 B.1: Problem Statement

2.1.1 Info on Hierarchical Time-Series Data:

Time series data usually refers to any data that changes value with respect to time. This could be the price of a commodity or stock, the demand for electricity, or the weather. Time series forecasting is widely used for the purpose of forecasting financial value, market demand, electricity demand, weather changes etc. For this reason, there is great interest in the research community to develop good forecasting algorithms for time series data.

One problem being investigated in regards to time series forecasting is making **coherent** predictions, i.e. predictions whose value adds up to a higher level of time series forecast. To understand this deeper, here is an explanation of **hierarchical time series data**.

Hierarchical time series refers to a set of different univariant time series data that are related to each other up the hierarchy through algebraic sums of one another. [1]

For example, if a time-series exist at the top-most hierarchy (level 0), and there are some more time-series existing at the immediate lower hierarchy (level 1), then the level 0 time series value at any point in time can be achieved by an algebraic sum of all its children time-series at its immediate lower level of hierarchy (level 0). A simple example is that for all sales data in every branch of a supermarket in a city, the sales data of all these stores must add up to the city-level sales data of that supermarket brand.




Figure 2-1A hierarchical time series level at level k along with its children series in the immediate lower level k+1 [2]

One of the interesting problems regarding hierarchical time series data, is while training machine learning algorithms to make a forecast on any time-series at a hierarchy, we can get better results by also including in the training, its children series at the immediate lower hierarchy. [2]

2.1.2 Problem Statement:

**Problem Statement:** In this project, I will train and evaluate the machine learning model based on architecture described in [2] to properly forecast hierarchical time-series data. I will use the dataset of Italian pasta sales used in [2] to train and evaluate this model.

2.1.3 Dataset Description:

**Dataset Name:** Hierarchical Sales Data

**Dataset website:** <https://archive.ics.uci.edu/dataset/611/hierarchical+sales+data>

**Dataset Download Website:** <https://data.mendeley.com/datasets/njdkntcpc9/1>

This dataset contains hierarchical sales data gathered from **an Italian grocery store**.

The dataset consists of 118 **daily time series** representing **the quantity of pasta products sold** from 01/01/2014 to 31/12/2018 of 4 national pasta brands. Besides univariate time series data, the quantity sold is integrated by information on the presence or the absence of a promotion. These time series can be naturally arranged to follow a 3-level hierarchical structure (see <https://www.sciencedirect.com/science/article/pii/S0957417421005431>). (Description copied from: [3])

**Dataset Characteristics = Time-Series**

**# Instances = 1798**

**# Features = 237**

As described in the previous section, this dataset has 237 feature columns including the date column. Hence there are 236 features to train on. Among these 118 features are **daily time-series data** of sales of pasta brands. The remaining 118 features are **explanatory variables/ external regressors** , i.e., further information about the sales like presence of daily promotion.

Hence,

237 features = 1 date column + 118 daily time series + 118 external regressors

Each time series has 5 years of sales data, with granularity of 1 day. (Some sales were recorded only in working days, hence not all days of these 5 years are present)

Below, is the dataframe shown with column names

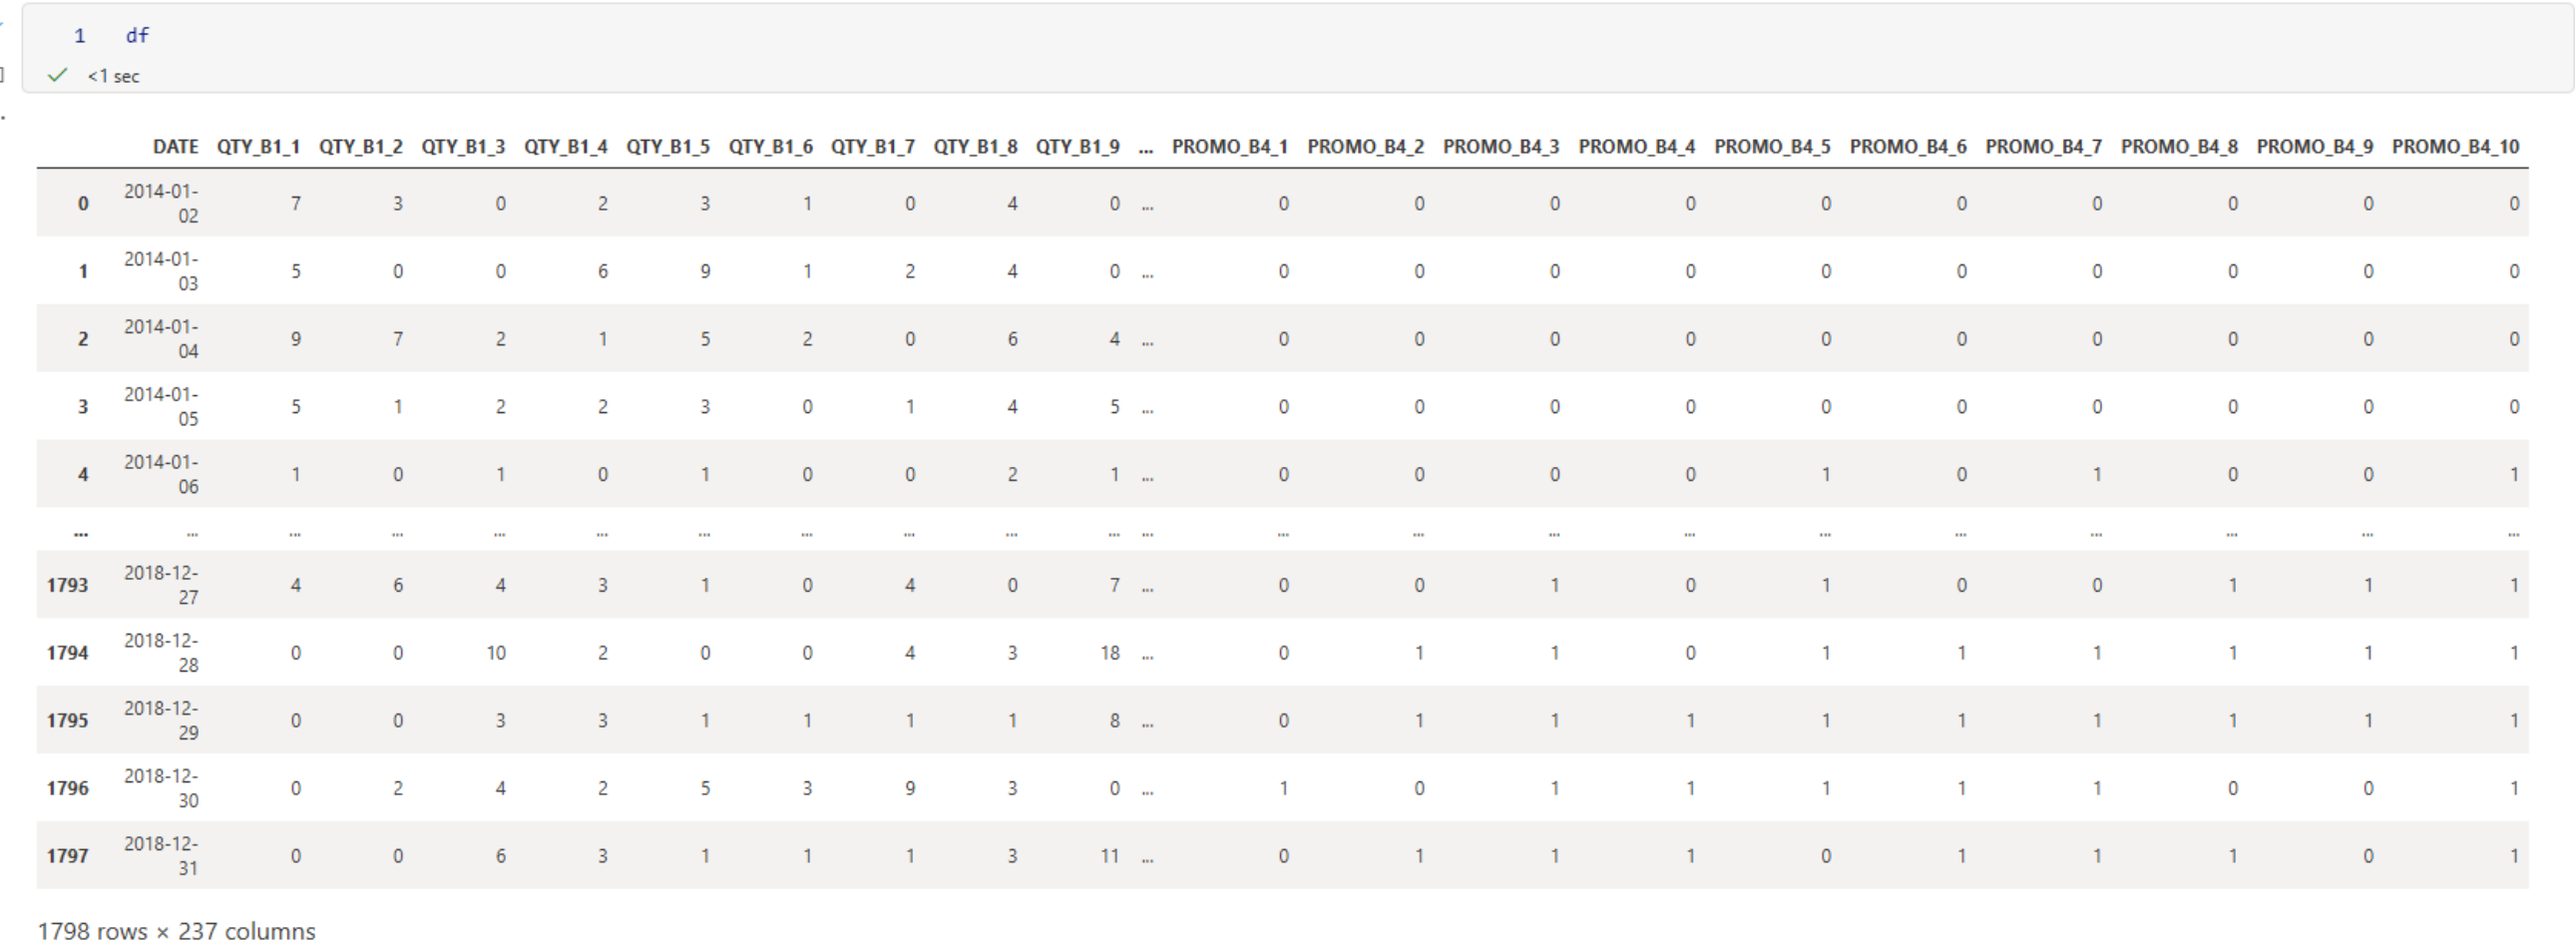


Figure 2-2 Code snippet showing the column names and values of the dataset

**Naming of Dataset Columns:**

- QTY\_B'X'\_'Y' - the quantity sold for brand 'X' item 'Y'
- PROMO\_B'X'\_'Y' - the promotion flag for brand 'X' and item 'Y')

**2.1.4 Modelling Plan:**

Summary of the data-preprocessing steps:

1. Identify the parent-child relations among the data columns, and group them together.
2. Concatenate the explanatory variables (promotion) data to the input time series data to create a single input vector
3. Create the 3-level hierarchical structure of the dataset.

- 4. Data-cleaning is unnecessary in this case as the dataset does not have missing values or irrelevant information.

The first step in data-preprocessing is to relate the relevant columns of the dataset to each other. The previous section outlines the different columns present in the dataset. Our first task is to group all the columns of pasta products (items Y) that fall under the same brand X. This will give us the parent-child relation of the hierarchical dataset.

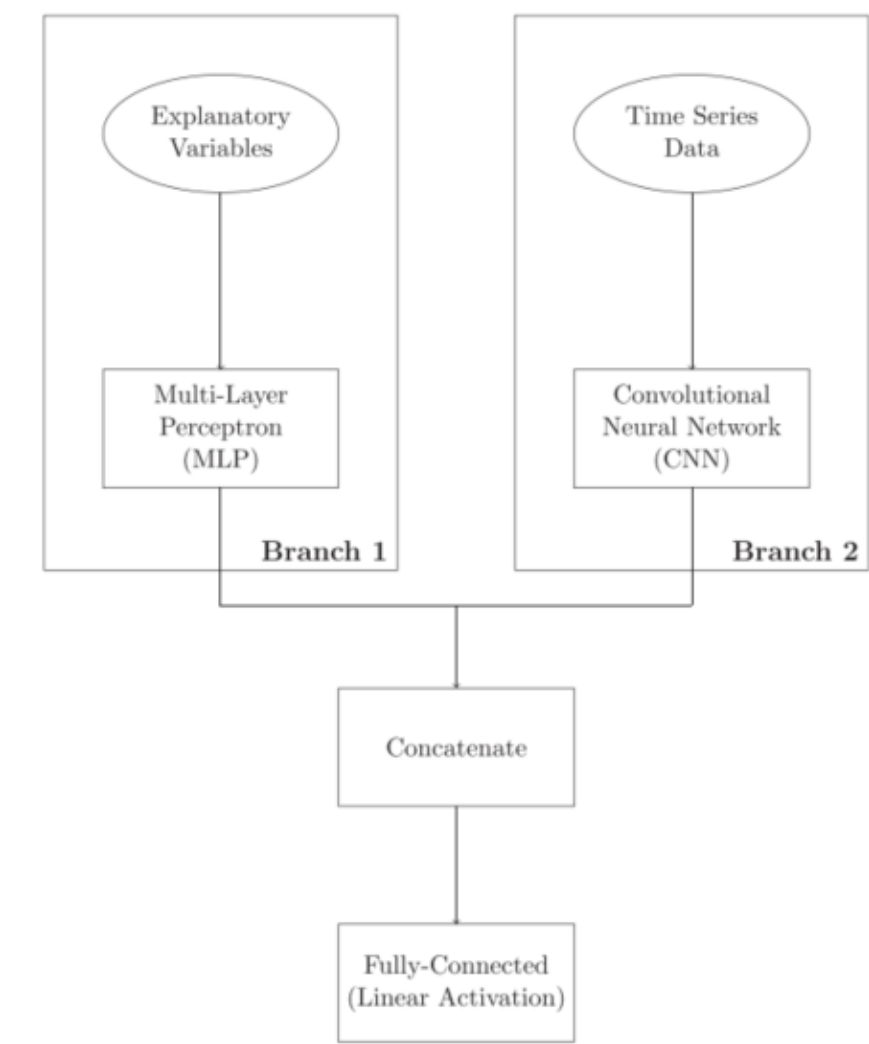


Figure 2-3The ML model will train on concatenated data, one part of the data is the time series data, another part is the explanatory variables (presence or absence of promotion). This input concatenation will be done separately for each distinct brand of pasta [2]

To train the ML model, we will concatenate the daily time series vector and explanatory variables vector to create a single input vector to train the ML forecasting model for each pasta brand. This concatenation of input vector will only contain the historic time series data of the series we are forecasting on, its immediate children in the hierarchy, and all explanatory variables (presence or absence of promotion of the included pasta brands or products). Hence for each different time series, the input vector will be different. We use the following naming convention described in the section – **Dataset Description**: to select the appropriate input vector for each time series forecast.

Lastly, we will create the hierarchical structure according to [2]

**2 level hierarchical structure implemented in [2]:**

- 1. Level-1: Sales data for each pasta brand (4 brands in total)

2. Level-2: Sales data for each item of the pasta brand. (Every item must fall into a child category of each of the above 4 mentioned branches.)

My hierarchy structure is different from [2] because I implement a 2-level hierarchy for simplicity whereas the paper [2] implements the 3 level hierarchy.

After doing all this data pre-processing and grouping of relevant data, we can move onto the next step of designing the ML architecture, and training the model


2.1.5 Data Pre-processing & Cleaning Requirements:

There might be some missing values in the dataset, these require to be filled by zero for proper time series predictions. Further description of the data cleaning and pre-processing work is mentioned in the following sections

2.2 Exploratory Data Analysis (EDA)

2.2.1 Daily Sales Time Series:

Here are some of the daily sales time series for some items across the entire time range (from 2014/01/01 to 2018/12/31)








1df.describe()  
✓ <1 sec



	QTY_B1_1	QTY_B1_2	QTY_B1_3	QTY_B1_4	QTY_B1_5	QTY_B1_6	QTY_B1_7	QTY_B1_8	QTY_B1_9	QTY_B1_10	...	PROMO_B4_1	PROMO_B4_2	PROMO_B4_3	PROMO_B4_4	I
count	1798.000000	1798.000000	1798.000000	1798.000000	1798.000000	1798.000000	1798.000000	1798.000000	1798.000000	1798.000000	...	1798.000000	1798.000000	1798.000000	1798.000000	
mean	5.895439	2.838154	2.994438	2.032814	2.694105	1.242492	2.567297	2.244160	5.066741	2.748610	...	0.204116	0.174082	0.193548	0.169633	
std	11.908115	3.844205	2.651103	1.733667	6.175587	1.539369	2.668616	2.211395	6.546864	8.116993	...	0.403166	0.379286	0.395189	0.375414	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	...	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	1.000000	1.000000	1.000000	0.000000	0.000000	1.000000	1.000000	1.000000	0.000000	...	0.000000	0.000000	0.000000	0.000000	
50%	3.000000	2.000000	2.000000	2.000000	1.000000	1.000000	2.000000	2.000000	3.000000	1.000000	...	0.000000	0.000000	0.000000	0.000000	
75%	6.000000	3.000000	4.000000	3.000000	3.000000	2.000000	3.000000	3.000000	6.000000	3.000000	...	0.000000	0.000000	0.000000	0.000000	
max	166.000000	41.000000	24.000000	12.000000	68.000000	12.000000	32.000000	17.000000	71.000000	246.000000	...	1.000000	1.000000	1.000000	1.000000	

8 rows × 236 columns

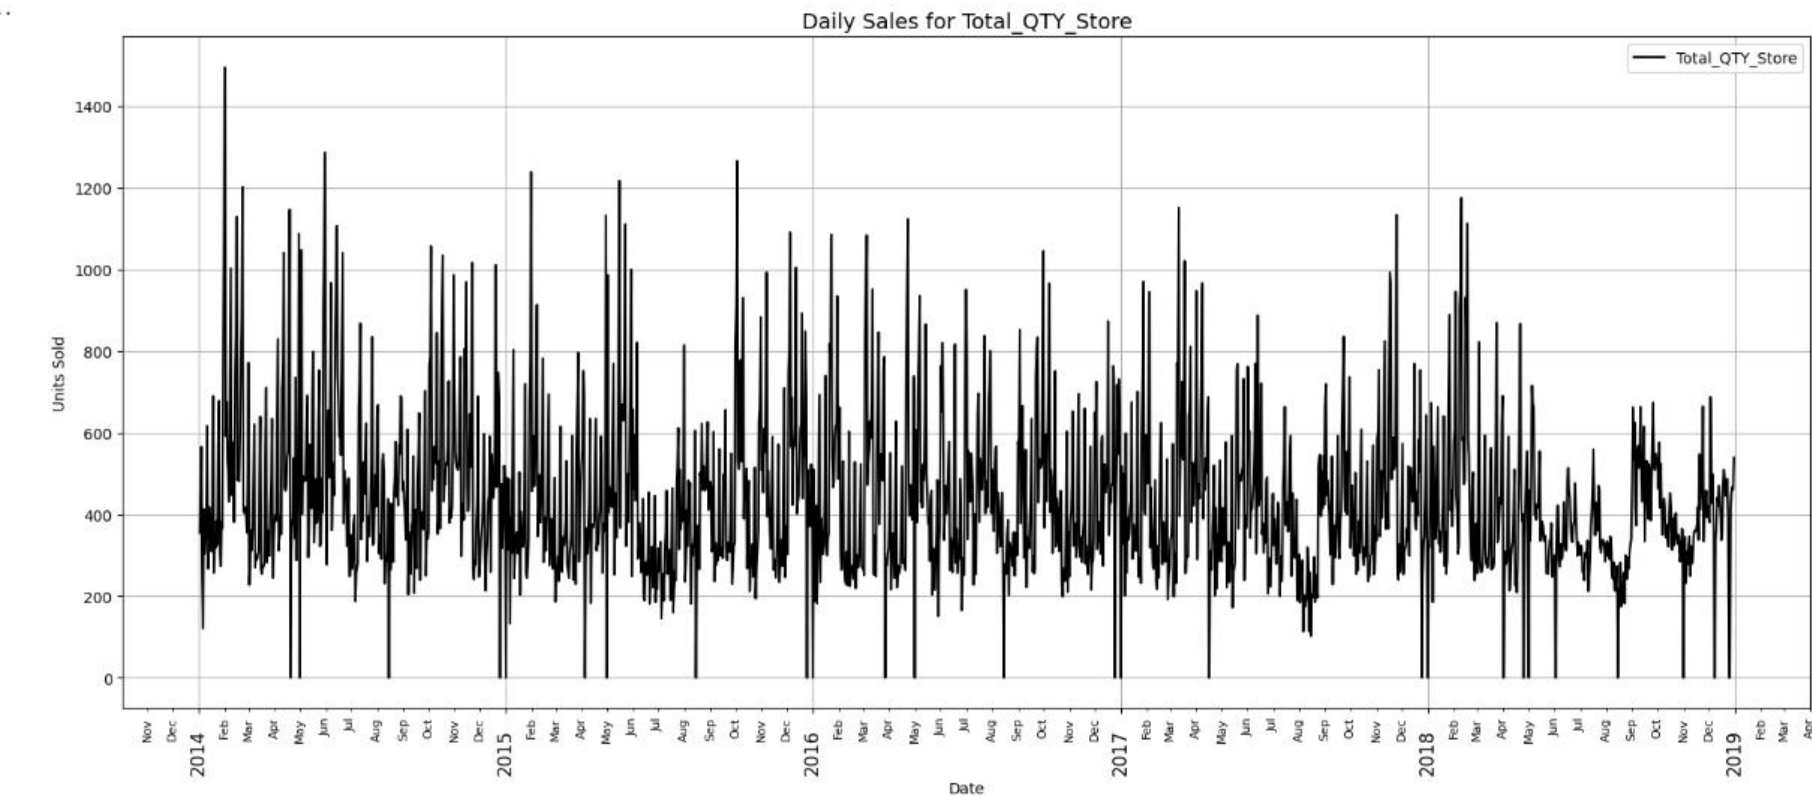


For all brands.



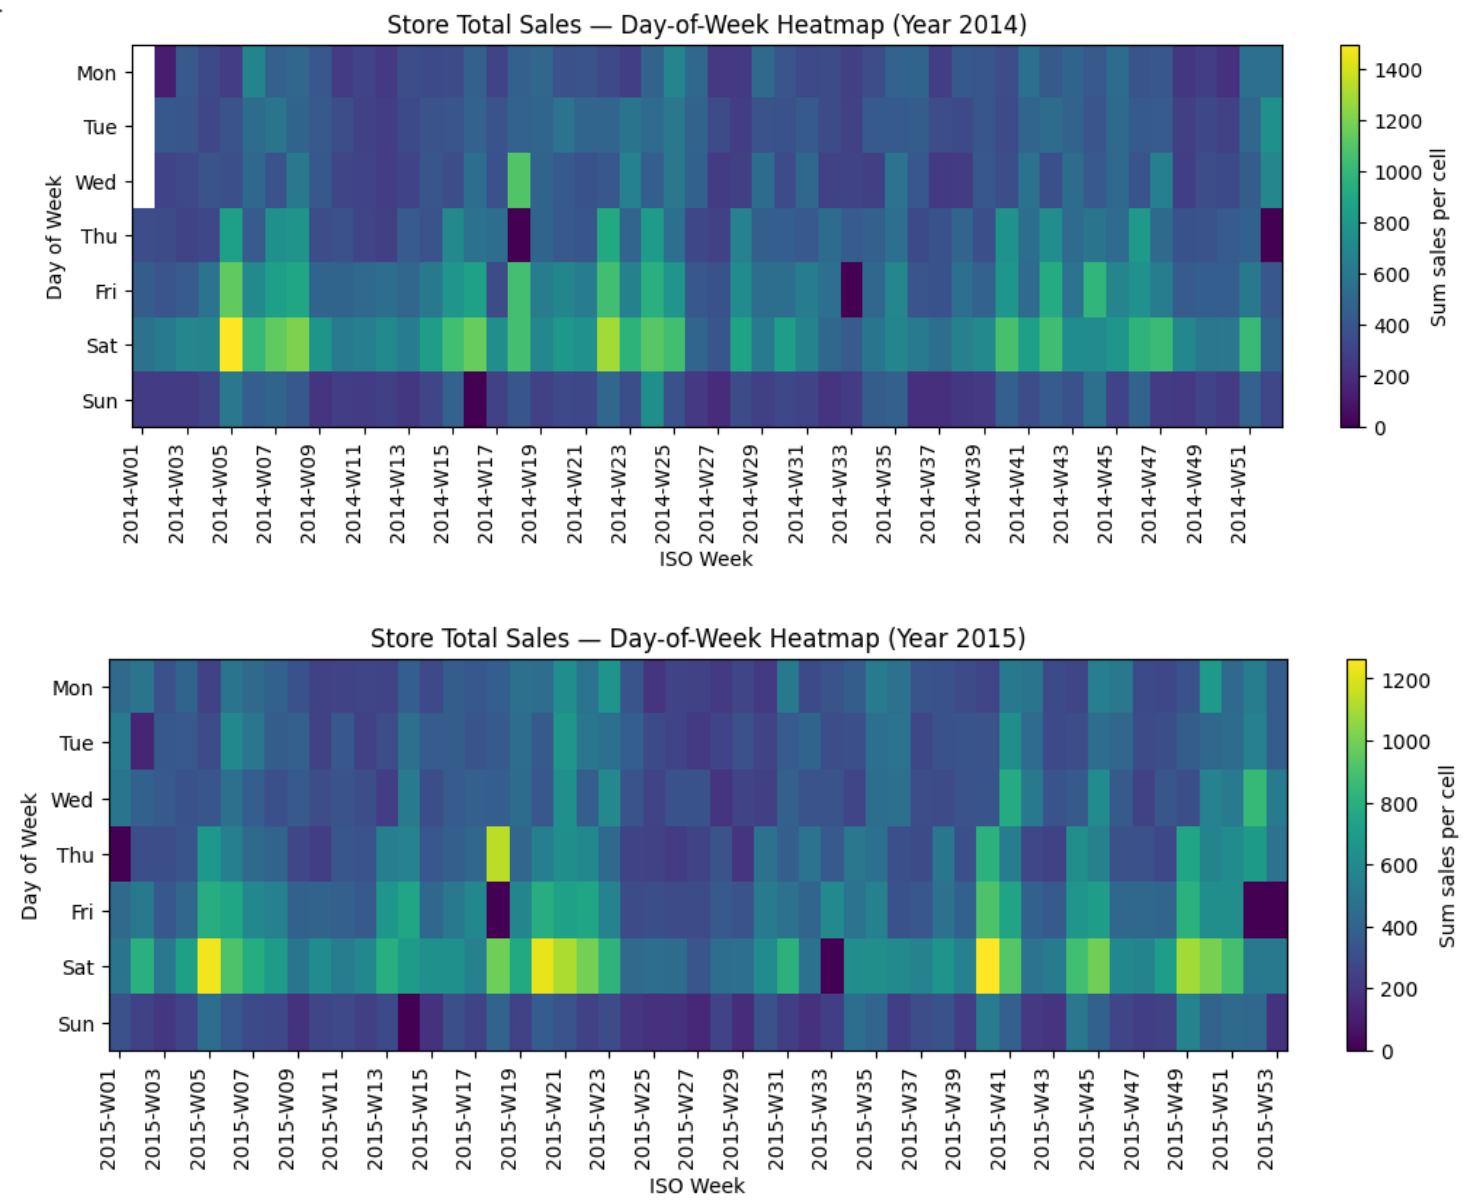



Total Sales:




2.2.2 Heatmap, Sales Across Weeks

The heatmaps represent the total sale each day of the week for each year. It helps identify recurring sales patterns in week days. The weeks were numbered using ISO system which starts number of each week from the first Thursday of that year.








The heatmaps clearly indicate that there are high sales in most Saturdays and some Fridays. Some portions of the heatmap is blank, indicating the dataset does not have sales data for some days which might be holidays when the store was closed.

2.2.3 Impact of Promotion


The plot shows the impact of promotion on top 20 items sold. The average daily sales were compared for this plot for each item. The top 20 displayed were selected on the basis of total sales in all 5 years, but the chart y-axis values are of average daily sales of each item.





2.2.4 Brand Sales Share by Year:


The pie charts describe the brand sales share in the store for each year





All the pie charts show that brand B1 has the major share in sales followed by brand B2. The sales shares are pretty consistent throughout the years.

2.2.5 Item Sales Share for each brand

Here are some plots showing the top 20 items sold from each brand, and stacked by year.





The plots obviously point out that some items of each brand have far greater amount in sales than others

2.3 Data Pre-processing & Cleaning:

The following steps of Data Cleaning and Pre-processing were done

2.3.1 Cleaning Data

- 1. Convert Time column to date-time
- 2. Clearing out missing values (store closed)
- 3. Filling those missing values with 0

2.3.2 Data Pre-process

- 1. Parsing to separate the brands and items
- 2. Defining functions to find sum between dates
- 3. (Augmenting the dataset) Adding new columns for brand level sales, total store sales
- 4. Scaling: Normalization on the data

2.4 ML Model 1:

The first ML model we trained was a simple **Ridge Regression Algorithm**. The theory of ridge regression is described in the following section.

To keep initial training simple, we applied this ML algorithm **only to Brand B1 sales** and **not in any other time series** of the dataset. Hence this ML model was trained without any relation to the other time series of this dataset, treating Brand B1 sales as an independent univariate time series data.

2.4.1 Ridge Regression: Theory

Theory of Ridge Regression

Ridge finds weights ( $w$ ) by minimizing the Objective function:

$$\underbrace{\sum_t (y_t - \hat{y}_t)^2}_{\text{fit error}} + \alpha \underbrace{\|w\|_2^2}_{\text{L2 penalty}} .$$

- ( $\alpha = 0$ )  $\rightarrow$  Ordinary Least Squares (no shrinkage; highest variance).
- **Larger** ( $\alpha$ )  $\rightarrow$  Stronger shrinkage; coefficients move toward 0 (lower variance, higher bias).

Note: the above picture is taken from the ipython notebook.

Here the  $y_t$  term is the actual series, and  $\hat{y}_t$  is the predicted time series. Alpha is a tuning parameter that we tune later

2.4.2 Feature Engineering

```
1 # Feature Matrix for training
2 X_all
✓ <1 sec
```

	lag_1	lag_7	lag_14	lag_28	roll_mean_7	roll_std_7	roll_mean_28	roll_std_28	dow_0	dow_1	dow_2	dow_3	dow_4	dow_5	dow_6	month_sin	month_cos
DATE																	
2014-01-30	120.0	111.0	100.0	101.0	143.857143	56.904682	136.464286	47.536717	False	False	False	True	False	False	False	5.000000e-01	0.866025
2014-01-31	144.0	181.0	133.0	136.0	148.571429	55.066194	138.000000	47.040566	False	False	False	False	True	False	False	5.000000e-01	0.866025
2014-02-01	181.0	256.0	250.0	162.0	148.571429	55.066194	139.607143	47.733317	False	False	False	False	False	True	False	8.660254e-01	0.500000
2014-02-02	302.0	113.0	97.0	106.0	155.142857	70.581935	144.607143	56.662967	False	False	False	False	False	False	True	8.660254e-01	0.500000
2014-02-03	97.0	91.0	124.0	47.0	152.857143	72.409549	144.285714	56.915049	True	False	False	False	False	False	False	8.660254e-01	0.500000
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
2018-12-27	0.0	187.0	145.0	130.0	138.857143	96.482172	132.500000	57.256732	False	False	False	True	False	False	False	-2.449294e-16	1.000000
2018-12-28	203.0	182.0	138.0	111.0	141.142857	97.990524	135.107143	58.780407	False	False	False	False	True	False	False	-2.449294e-16	1.000000
2018-12-29	192.0	204.0	138.0	139.0	142.571429	98.755349	138.000000	59.538347	False	False	False	False	False	True	False	-2.449294e-16	1.000000
2018-12-30	158.0	227.0	131.0	134.0	136.000000	95.462034	138.678571	59.658319	False	False	False	False	False	False	True	-2.449294e-16	1.000000
2018-12-31	182.0	172.0	122.0	147.0	129.571429	89.650910	140.392857	60.206029	True	False	False	False	False	False	False	-2.449294e-16	1.000000

1797 rows × 17 columns

Figure 2-4 Feature Matrix for Ridge Regression

The above picture shows the feature matrix for the input to the ML algorithm for ridge regression. Weights will be multiplied to all the features of this feature matrix, and then the optimal weights will be found during training of the ridge regression algorithm.

2.4.2.1.1 Lag Features:

The Figure 2-4 terms lag\_1, lag\_7, lag\_14, lag\_28, each are respectively the sales values of 1 day before, 7 days before, 14 days before, and 28 days before the selected date in the row. Hence the number of rows has also reduced by 28 from 1825 to 1797 as seen in the above picture. This is to ensure that lag does not cause going behind to a value outside of the time range of the dataset. Hence a first few values have been omitted.

These lag features help look at the recent history of those days before and feed those features into the learning algorithm.

2.4.2.1.2 Rolling Statistics

Similarly, the **Rolling statistics** roll\_mean\_7, roll\_std\_7, roll\_mean\_28, roll\_std\_28, are the mean and standard deviation for a rolling window of 7 days and 28 days behind the current point. This window statistics is also fed to the ridge regression algorithm for learning.

### 2.4.2.1.3 Days of Week (Dow) Encoding

Days of the week (DOW) has been encoded in a **one-hot-encoding** format as shown from the Figure 2-4 Feature Matrix for Ridge Regression

### 2.4.2.1.4 Months of Year, cyclic encoding

Months of the year has been encoded in **cyclic encoding** to keep January, close to December. The reason this was done was because of the seasonal patterns similarity existing in January, December (cold months) vs June, July (hot months).

Whereas one-hot encoding was done for DOW because there might be sudden changes in customer behavior from Sunday to Monday, hence not ideal for cyclic encoding.

## 2.4.3 Train-Test Split

Next we did a train-test split in the dataset according to the following figure

### Split: Time-aware split (train/test/vald) by horizon lengths

```

1  # -----
2  # 3) Time-aware split (train/valid/test) by horizon lengths
3  # -----
4  test_horizon = 365 # last 365 days for final evaluation
5  valid_horizon = 365 # preceding 365 days for model selection
6
7  n = len(y_all)
8  if n < (test_horizon + valid_horizon + 30):
9      # Ensure we have enough history; shrink horizons if needed
10     valid_horizon = max(60, min(valid_horizon, n // 5))
11     test_horizon = max(60, min(test_horizon, n // 5))
12
13  test_start = n - test_horizon
14  valid_start = test_start - valid_horizon
15
16  X_train, y_train = X_all.iloc[:valid_start], y_all.iloc[:valid_start]
17  X_valid, y_valid = X_all.iloc[valid_start:test_start], y_all.iloc[valid_start:test_start]
18  X_test, y_test = X_all.iloc[test_start:], y_all.iloc[test_start:]
19
20  print("Train range:", X_train.index.min().date(), "→", X_train.index.max().date(), f"({len(X_train)} days)")
21  print("Valid range:", X_valid.index.min().date(), "→", X_valid.index.max().date(), f"({len(X_valid)} days)")
22  print("Test range:", X_test.index.min().date(), "→", X_test.index.max().date(), f"({len(X_test)} days)")

```

✓ <1 sec

Train range: 2014-01-30 → 2016-12-31 (1067 days)  
Valid range: 2017-01-01 → 2017-12-31 (365 days)  
Test range: 2018-01-01 → 2018-12-31 (365 days)

Figure 2-5 Train-test split by horizon length

We specified the test-horizon to be 365 day, and the validation horizon to be 365 days. That is it will take the last 365 + 365 days for test and validation and will use the remaining days to train the ridge regression model, seen from Figure 2-5

2.4.4 Training & Cross-Validation

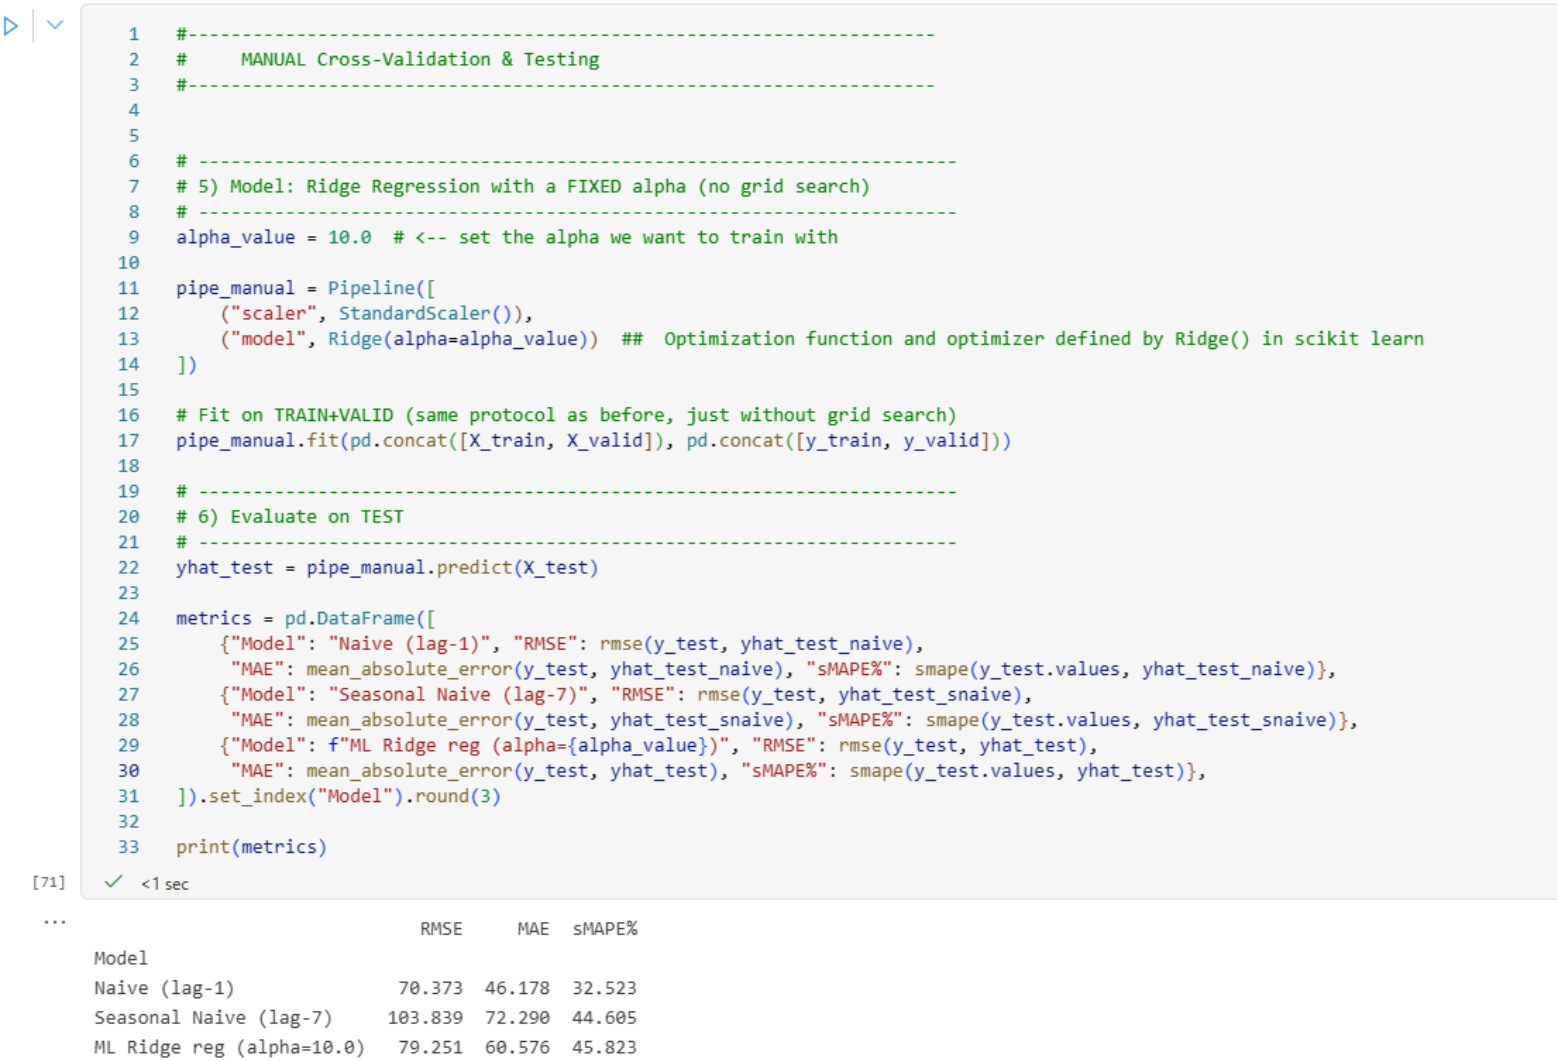


Figure 2-6 Training Fit to Ridge Regression

The above Figure 2-6 shows the training of the ridge regression model. The **ridge regression implementation from scikit-learn** is used in this training. This implementation uses an auto solver that uses the Cholesky decomposition to solve for optimal weights.

After training, the results are shown in Figure 2-6.

Results used are Root Mean Squared (RMSE), Mean Absolute Error (MAE), Symmetric Mean Absolute Percentage Error (sMAPE)

The ML Ridge regression model is compared with some other prediction methods, the **Naïve**, and **Seasonal Naïve** . Naïve and Seasonal Naïve does no calculations at all, Naïve just draws the lag\_1 value (the immediate preceeding value) from the feature matrix, and Seasonal Naïve draws out the lag\_7 value (the value 7 days behind this value) from the feature matrix. They are just used as markers to see how the ML model compares to these easy predictors.

2.4.5 Hyperparameter Tuning

Hyperparameter tuning was later done on this model using grid search.

Hyperparamter Tuning ML 1  
GridSearch Cross Validation

3]

✓ 2 sec

```
1 # -----
2 # 5) Model: Ridge Regression with TimeSeriesSplit and grid search on alpha
3 # -----
4 pipe = Pipeline([("scaler", StandardScaler()), ("model", Ridge())])
5
6 tscv = TimeSeriesSplit(n_splits=16) # Splits into 4 folds of validation
7
8 # Cross Validation Parameter Grid
9 # Alpha
10 param_grid = {"model__alpha": [0.1, 1.0, 3.0, 10.0, 30.0, 100.0, 120.0, 140.0, 180.0, 200.0, 220.0, 260.0, 280.0, 300.0, 320.0, 330.0, 350.0, 400.0]}
11 # Alpha is the L2 regularization parameter
12 # It helps prevent overfitting by penalizing large coefficients, allowing generalization in prediction, better performance in test
13
14
15 # GridSearch Cross-Validation
16 gcv = GridSearchCV(pipe, param_grid, scoring="neg_mean_squared_error", cv=tscv)
17 # RMSE metric used for Cross Validation, to pick the best model
18 # -ve RSME means better performance the closer to 0 (absolutes of larger negative numbers are smaller)
19
20 # Other cv scoring metrics
21 # scoring = "neg_mean_absolute_error"
22 # scoring = "neg_mean_absolute_percentage_error"
23 # scoring = "neg_mean_squared_error"
24 # scoring= "neg_root_mean_squared_error"
25
26
27 gcv.fit(X_train, y_train)
28
29 best_model = gcv.best_estimator_
30
31 # Refit on TRAIN+VALID
32 print("Best Model: from Grid Search")
33 best_model.fit(pd.concat([X_train, X_valid]), pd.concat([y_train, y_valid]))
```

Best Model: from Grid Search

Pipeline

StandardScaler

Ridge

Ridge(alpha=30.0)

Figure 2-7 Hyperparameter tuning for ML Ridge Model



The hyperparameter tuning was done on a grid search. The hyperparamter tuned here was the \alpha parameter.

The alpha was the L2 regularization parameter shown in Figure 2-8

# Explanation of Alpha Parameter

$\alpha$  (Alpha): We Tune Alpha

**Alpha** is the **L2-regularization strength** in Ridge regression — a knob that controls how much we shrink the model’s coefficients to combat overfitting.

## Ridge Regression: Objective Function (and role of Alpha)

Ridge finds weights ( $w$ ) by minimizing:

$$\underbrace{\sum_t (y_t - \hat{y}_t)^2}_{\text{fit error}} + \alpha \underbrace{\|w\|_2^2}_{\text{L2 penalty}}.$$

- ( $\alpha = 0$ ) → Ordinary Least Squares (no shrinkage; highest variance).
- **Larger** ( $\alpha$ ) → Stronger shrinkage; coefficients move toward 0 (lower variance, higher bias).
- In scikit-learn’s `Ridge`, the **intercept is not penalized**; only feature weights are.

## We Grid-Search Alpha

We don’t know beforehand how much shrinkage gives the best **out-of-sample** performance.  
So we try a set of values (e.g., `0.1, 1, 3, 10, 30, 100`) and use **time-aware cross-validation** to pick the ( $\alpha$ ) that forecasts best on later, unseen periods.

Figure 2-8 Ridge Regression Alpha Paramter Tuning

Higher values of alpha allows for generalization.

After implementing grid search we found the best model had alpha = 30.0 as shown in Figure 2-7

2.4.6 Results from Best model for ML model 1

Results- Best From Cross Validation



Figure 2-9 Best Results from Cross Validation of ML model 1




Figure 2-10 Plot of B1 Brands total store sales, actual and predicted

The forecast results for the trained ridge model is shown in Figure 2-10

2.5 ML Model 2:

2.5.1 Architecture for Hierarchy Exploitation:

This ML Model no. 2 exploits the hierarchical relationship between brand and its items. Hence the ML model will train together the time series for the brand and time series for every item in the brand as shown in Figure 2-1

This ML model fully uses the modelling plan described in section Modelling Plan:

The Time series hierarchy and Modelling plan is also presented below.




Figure 2-11Time Series Hierarchy followed for ML Model 2




Figure 2-12 Neural Network Disaggregation (NND) architecture followed in paper [2]

The Neural Network Dissaggregation (NND) architecture shown in Figure 2-12 and adapted from [2] is implemented in this ML Model 2.

But one difference from this model described in [2] and my implementation is that **I only used hierarchy of brand level to item level**. The **final hierarchy of total store level sales was not used** to keep the code simple.

2.5.2 Feature Engineering

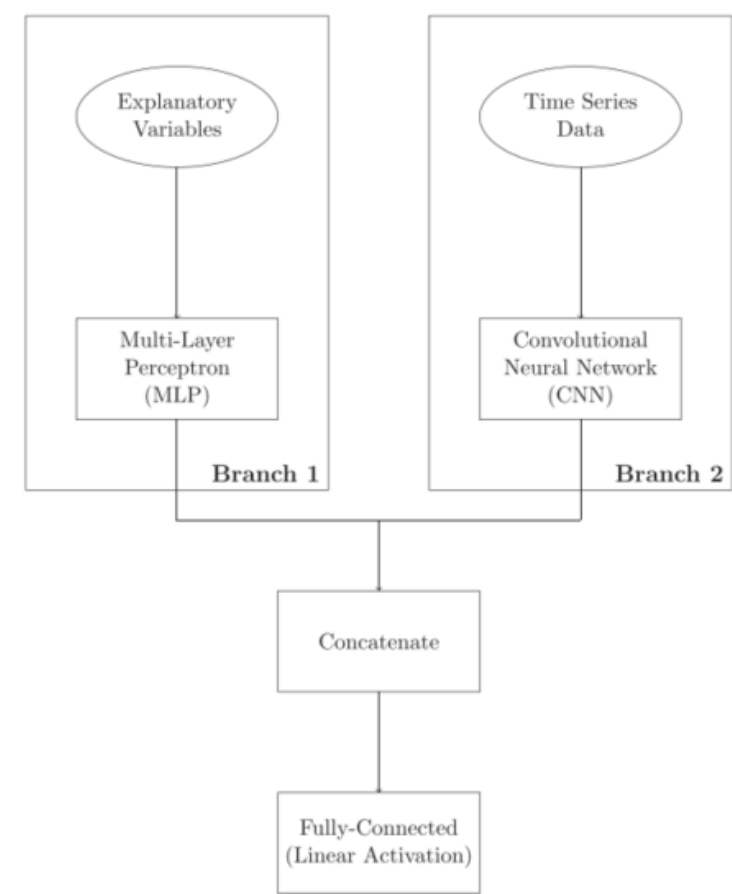


Figure 2-13 Input Variable Concatenation plan for ML Model 2. Adapted from [2]

For this NND architecture to work and train on both the top level time series and its contributing children time series, we have to concatenate the training data (previous historical data of the time series) and the **exogenous / explanatory variables** (in this case, the presence or absence of promotion)

These exogeneous variables were extracted from every brand and every item using the code shown in Figure 2-14

Meanwhile Figure 2-2 shows the original dataset which contains the exogenous variables for each brand. They were separated using the code from Figure 2-14

## Exogenous Variable Extraction for Brand

```

1  # -----
2  # 1) Utilities to extract brand & item series and exogenous features
3  # -----
4
5  from typing import Dict, List
6
7  def get_items_for_brand(brand: str) -> List[str]:
8      items = brand_dict.get(brand, [])
9      if not items:
10         raise ValueError(f"No items found for brand {brand} in brand_dict.")
11     return items
12
13
14 def get_brand_series(frame: pd.DataFrame, brand: str) -> pd.Series:
15     # If a precomputed brand column exists, prefer it
16     col = f"Brand_QTY_{brand}"
17     if col in frame.columns:
18         y = frame[col].astype(float)
19     else:
20         # Sum item-level quantities
21         items = get_items_for_brand(brand)
22         qty_cols = [f"QTY_{brand}_{i}" for i in items if f"QTY_{brand}_{i}" in frame.columns]
23         if not qty_cols:
24             raise ValueError(f"No QTY columns for {brand} present in df.")
25         y = frame[qty_cols].sum(axis=1).astype(float)
26     return y.asfreq("D").fillna(0.0)
27
28
29 def build_brand_exog(frame: pd.DataFrame, brand: str) -> pd.DataFrame:
30     ex = pd.DataFrame(index=frame.index)
31     # Promo share = relative number of items under promo for this brand
32     items = get_items_for_brand(brand)
33     promo_cols = [f"PROMO_{brand}_{i}" for i in items if f"PROMO_{brand}_{i}" in frame.columns]
34     if promo_cols:
35         ex["promo_share"] = frame[promo_cols].sum(axis=1) / float(len(promo_cols))
36     else:
37         ex["promo_share"] = 0.0
38     # Calendar dummies: DOW & Month (drop_first to avoid dummy trap if intercept exists)
39     dow = pd.get_dummies(frame.index.dayofweek, prefix="dow", drop_first=True)
40     mon = pd.get_dummies(frame.index.month, prefix="mon", drop_first=True)
41     dow.index = ex.index; mon.index = ex.index
42     ex = pd.concat([ex, dow, mon], axis=1).astype(float)
43     return ex.asfreq("D").fillna(0.0)
44
45
46 def build_item_exog(frame: pd.DataFrame, brand: str, items: List[str]) -> Dict[str, pd.DataFrame]:
47     # Shared calendar dummies
48     dow = pd.get_dummies(frame.index.dayofweek, prefix="dow", drop_first=True)
49     mon = pd.get_dummies(frame.index.month, prefix="mon", drop_first=True)
50     dow.index = frame.index; mon.index = frame.index
51     calendar = pd.concat([dow, mon], axis=1).astype(float)
52
53     exog = {}
54     for it in items:
55         promo_col = f"PROMO_{brand}_{it}"
56         if promo_col in frame.columns:
57             promo = frame[promo_col].astype(float)
58         else:
59             promo = pd.DataFrame({promo_col: np.zeros(len(frame), dtype=float)}, index=frame.index)
60         exog[it] = pd.concat([promo, calendar], axis=1).asfreq("D").fillna(0.0)
61     return exog
62
63
64 def get_items_matrix(frame: pd.DataFrame, brand: str, items: List[str]) -> pd.DataFrame:
65     cols = [f"QTY_{brand}_{it}" for it in items if f"QTY_{brand}_{it}" in frame.columns]
66     if len(cols) != len(items):
67         missing = [it for it in items if f"QTY_{brand}_{it}" not in frame.columns]
68         raise ValueError(f"Missing QTY columns for items: {missing}")
69     return frame[cols].astype(float).asfreq("D").fillna(0.0)

```

Figure 2-14 Code for Exogenous variable extraction of both items and brands

2.5.3 Code for ML Architecture

```
1 # -----
2 # 3) Models
3 # -----
4
5 def build_brand_cnn_mlp(window: int, n_exog: int) -> Model:
6     # y branch
7     inp_y = layers.Input(shape=(window, 1), name="y_seq")
8     a = layers.Conv1D(16, 3, padding="causal", activation="relu")(inp_y)
9     a = layers.Conv1D(16, 3, padding="causal", activation="relu")(a)
10    a = layers.GlobalAveragePooling1D()(a)
11    # X branch
12    inp_x = layers.Input(shape=(window, n_exog), name="x_seq")
13    b = layers.Flatten()(inp_x)
14    b = layers.Dense(64, activation="relu")(b)
15    # fuse
16    h = layers.Concatenate()([a, b])
17    h = layers.Dense(64, activation="relu")(h)
18    out = layers.Dense(1, name="yhat_scaled")(h) # trained in scaled space
19    model = Model([inp_y, inp_x], out)
20    model.compile(optimizer="adam", loss="mse")
21    return model
```

[80] ✓ <1 sec

▶ | ▾

```
1 # Disaggregation model architecture
2
3 # NDD
4
5 def build_disagg_nnd(window: int, n_items: int, n_item_exog: int,
6     brand_latent_dim: int = 16) -> Model:
7     """NND that outputs per-item forecasts coherent with the brand scalar.
8     Inputs:
9     - yB_seq: (w,1) brand history
10    - Xitem_seq: (I,w,m) per-item exog sequence
11    - yB_scalar: (1,) brand scalar at target time (true for training, forecast at inference)
12    Output:
13    - Yitem_pred: (I,) per-item forecasts that sum to yB_scalar (softmax shares * scalar)
14    """
15    # Brand branch
16    inp_yB = layers.Input(shape=(window, 1), name="yB_seq")
17    a = layers.Conv1D(brand_latent_dim, 3, padding="causal", activation="relu")(inp_yB)
18    a = layers.GlobalAveragePooling1D()(a) # (batch, brand_latent_dim)
19
20    # Per-item exog branch
21    inp_Xi = layers.Input(shape=(n_items, window, n_item_exog), name="Xitem_seq")
22    # Flatten time+features for each item, keep item axis
23    z = layers.Reshape((n_items, window * n_item_exog))(inp_Xi)
24    z = layers.TimeDistributed(layers.Dense(64, activation="relu"))(z)
25    z = layers.TimeDistributed(layers.Dense(32, activation="relu"))(z) # (batch, I, 32)
26
27    # Repeat brand latent across items and fuse
28    a_rep = layers.RepeatVector(n_items)(a) # (batch, I, brand_latent_dim)
29    u = layers.Concatenate(axis=-1)([z, a_rep]) # (batch, I, 32+latent)
30    u = layers.TimeDistributed(layers.Dense(32, activation="relu"))(u)
31    logits = layers.TimeDistributed(layers.Dense(1))(u) # (batch, I, 1)
32    logits = layers.Reshape((n_items,))(logits) # (batch, I)
33
34    # This part ensures coherence as sums to 1 for the brand level
35    shares = layers.Softmax(axis=-1, name="shares")(logits) # sums to 1 across items
36
37    # Multiply by brand scalar to get coherent item forecasts
38    inp_yB_scalar = layers.Input(shape=(1,), name="yB_scalar") # (batch,1)
39    yB_rep = layers.Concatenate()([inp_yB_scalar for _ in range(n_items)]) # (batch, I)
40    y_items_pred = layers.Multiply(name="Yitem_pred")(shares, yB_rep)
41
42    model = Model([inp_yB, inp_Xi, inp_yB_scalar], y_items_pred)
43    model.compile(optimizer="adam", loss="mse") # MSE on items in original units
44    return model
```

[81] ✓ <1 sec

Figure 2-15 Code for NDD architecture

The code for the NDD architecture is shown in Figure 2-15. Here, in the cell # 2, and line # 34, we can see the enforcement of the **coherence constraint** (ie, sum of all child time series must equal its parent time series predictions). This ensures that the model learns relations such that sum of all its child series will equal to the prediction in the parent series. This helps the model learn better.

## 2.6 Results of ML Model 2

After training, the following results were obtained for the ML model 2 using the NND.



Figure 2-16 Results, NND on B1 sales






Figure 2-17 B1 Sales Plot

The above figure show the forecast made by NND model after training. The forecast well matches the actual data. This is for B1 sales only.

2.7 Hyperparameter Tuning for ML Model 2 (NND)

Hyperparameter tuning was done on the B1 sales results to simplify the search.

Hyperparameter tuning was done on the following features:

- 1. Window size
- 2. Brand Latent Dimension

2.7.1 Window Size Tuning

Window size determines the window size to look for during the training process. Hence we tried three different window sizes, 10, 30, 60. Results are shown below.

Window size	Root Mean Squared Error (RMSE)	Mean Absolute Error (MAE)
10	85.252	61.74
30	88.382	61.856
60	93.346	68.156

Clearly Window =10 was the best performer. So we selected window = 10 for following calculations

2.7.2 Brand Latent Dimension Tuning



2.7.3 Best Hyperparamters

Best Hyperparameters

Window = 10

Brand Latent Dimension = 16

2.8 Results: All Brands

Using the best hyperparameter, the following results of forecast for all remaining brands are shown




Figure 2-19 Brand B2 Forecast




Figure 2-20 Brand B3 Forecast




Figure 2-21Brand B4 Forecast

As we can see, all the figures show that the forecast is pretty accurate

2.9 Conclusions:

2.9.1 Comparing ML Model 1 and ML Model 2

Comparing ML model 1 (Ridge Regression with no information of hierarchy) vs. ML Model 2 (NND Model with hierarchy information) we can clearly see that the forecast of ML model 2 (Figure 2-17, Figure 2-19, Figure 2-20, Figure 2-21) are much better. The figures visually show better forecasts then the forecast of ML model 1 (Figure 2-10)

Hence the NND model was able to generate better forecasts by including information of the hierarchy in the brands to its items.

2.9.2 Future Work Scope

Although in this project, I was not able to use the full level of hierarchy, (I omitted the top level of store level sales that accumulate all brand sales) it can be left to future work to implement this 3 level hierarchy in the NND architecture and train all 3 hierarchies of time series simultaneously.

### 3 References

[1] R. & A. G. Hyndman, "Forecasting: principles and practice.," OTexts, 2018.

[2] P. M. a. V. P. a. A. M. Sudoso, "A machine learning approach for forecasting hierarchical time series," *Expert Systems with Applications*, p. 115102, 2021.

[3] P. M. a. V. P. a. A. M. Sudoso, "UC Irvine Machine Learning Repository," [Online]. Available: <https://archive.ics.uci.edu/dataset/611/hierarchical+sales+data>.